

Ansible, best practices.

Bas Meijer 2016 Amsterdam

Download the project

```
brew install python ansible  
git clone https://github.com/bbaassssiiee/lunchbox  
vagrant up
```

- Requires: VirtualBox, Vagrant

Topics in this workshop

- Ansible.
- Get organized.
- Format with style.
- Inventories and environments.
- Variable precedence.
- A vault with secrets.
- Build quality in.
- Roles in detail.
- Managing efforts.
- Benchmarking where you are.

What is Ansible?

- A radically simple IT automation engine for:
 - cloud provisioning
 - configuration management
 - application construction and deployment
 - intra-service orchestration
 - many other IT needs

Why Ansible?

- Simple to learn and share. Easy to audit.
- Global open source community.
- Minimal requirements.
- Secure, only SSH.
- Integrates really well with cloud, Docker, etc.
- Role-based access & delegation in Tower.

Application Construction

- vagrant, local mode
- packer, image creation
- docker build
- cloud & bare metal provisioning
- deployment & remote execution

Get organized

- Try to schedule provisioning and deployment as early and often as possible.
- Roll-out changes, in stages, to all environments using a CI server and Tower.
- Use the same playbooks in each environment.
- Create playground environments to test things.

Be Idempotent

- Same operation? Same result, again & again.
- Ensure no changes unless things change.
- No uncertainty: describe desired state.
- System life cycle as transaction log, accounts for all changes done on purpose.

Use version control

- Collaboration and preservation.
- Local validation, merge-requests, tests.
- git-flow: audit trail / 4 eyes.
- Role sharing / joy of modular playbooks.
- Galaxy & GitHub have a wealth of stuff.

Git version control

- `git init`
- `git add playbook.yml inventory`
- `git commit -m "why I made changes" -a`
- `git push`, after you tested those changes
- *`echo ansible.cfg >> .gitignore`*

Keep it simple

- If you can make it simple, make it simple.
- If something feels complicated, it probably is.
- Always define state.
- Do one thing at a time, commit related changes.
- Make it readable for your sucesor.

Organize your content

- Format YML: with style.
- Magic happens when you put your files in the right locations.
- Editors use file extensions for syntax coloring.
- Version control and Ansible go hand in hand.

Whitespace and comments

- # comments start with a hash-mark.
- Yml indents with 2 spaces.
- White space helps readability.
- Add a blank line before vars, roles, tasks, handlers.

Directory layout

- `group_vars`
- `host_vars`
- `roles`

```
ansible.cfg      # parameters that affect running ansible
vagrant.ini      # inventory file define the hosts in an environment
group_vars/    # here we assign variables to particular groups
    all          # global variables for all groups
    dbservers/   # directory for dbservers group
        secrets  # -- encrypted variables for dbservers group
        vars     # -- plaintext variables for dbservers group
    group2       # plaintext variables for group2
host_vars/     # here we assign variables to particular hosts
    hostname1    # if systems need specific variables, put them here
    hostname2    # ""
vagrant.yml      # master playbook for environment 'vagrant'
webservers.yml   # playbook for webserver tier
dbservers.yml    # playbooke for database tier
galaxy_roles/  # roles imported from galaxy
roles/        # in-house roles
    common/      # this hierarchy represents a "role"
        tasks/   # 'tasks' contains the actions that implement role
            main.yml # -- main.yml could include other files if warranted
        handlers/ # 'handlers' can be notified by tasks on change
            main.yml # -- handlers file often defines service actions
        templates/ # files for use with the template module
            hosts.j2 # -- Jinja templates, should end in .j2
        files/     # 'files' is the start for relative paths
            static.txt # -- file for use with the copy module
            script.sh # -- script for use with the script module
```

Format your playbooks

- shebang #!
- use a name
- tags in header
 - preparation
- blank lines
- linebreak=80
- align args

```
#!/usr/bin/env ansible-playbook

- name: 'install.yml'           # quote names for syntax highlighting
  hosts: localhost              # scope the play appropriately
  connection: local            #
  gather_facts: False          # booleans: /^(y|yes|n|no|true|false|on|off)$/i

  tags:                         # use tags for plays, and actions
    - preparation

  vars:                         # use group_vars for environment specifics
    - url: "https://galaxy.ansible.com" # quote when value has ':'

  tasks:                       # list tasks, but consider using a role
    - name: 'check network'     # format parameters for small terminal size
      uri: url={{ url }}
          method=HEAD
          return_content=no
          status_code=200
          timeout=60
          follow_redirects=all

    - name: 're-import roles from Galaxy'
      command: ansible-galaxy install --force --role-file requirements.yml
```


Tag all the things

- Tags help organizing playbooks.
- Tags can help in testing.
- You can run or skip parts of playbooks:
 - `--tags=only,run,these,tags`
 - `--skip-tags=tags,to,skip`

Deployment vs configuration

- Deploy all?
- --tags
- --limit
- Ad-Hoc
- playbooks/

```
ansible-playbook -i production site.yml -vv
```

```
ansible-playbook -i production webservers.yml
```

```
ansible-playbook -i production site.yml --tags ntp
```

```
ansible-playbook -i acceptance dbservers.yml --tags restore
```

Start your model with the inventory

- Hosts used in infrastructure/cloud
- Arbitrary grouping
- Can be used for staging (vagrant, int, acc, prod).

```
[local]  
127.0.0.1 ansible_connection=local
```

```
[vagrant:children]  
dbservers  
webservers
```

```
[dbservers]  
sql
```

```
[webservers]  
web
```

Inventories

- Use .ini files for static inventories.
- Executable inventory file must emit JSON.
- If you have a system maintaining a list of systems in your infrastructure, use it for your dynamic inventory.
- Ansible Tower keeps inventories and credentials securely.

Stage your environments

- Use separate inventories for stages like 'test' and 'production'.
- Target an environment with the -i flag.
- Test things first in a stage environment before running them in 'production'.
- Use group_vars to manage differences.

Variables

- All systems are not exactly alike, use variables to deal with differences.
- Separate environment/staging data from configuration.
- Use the template language Jinja2.

```
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1         localhost localhost.localdomain localhost6 localhost6.localdomain6

{% for item in groups['all'] %}
{{ hostvars[item]['ansible_ssh_host'] }} {{ item }}
{% endfor %}
```

Group variables

- Hosts are listed in the inventory under [groups]
- The group_vars directory has the configuration data in files and folders.
- Use a subdir for each group with a 'vars' and 'secrets' file.

Use vars to define your values

- Roles have defaults.
- `group_vars` is for overrides.
- 'all' is for globals

```
---
# file: group_vars/all
# here we assign variables to all groups

ansible_user: ansible

google_nameserver: 8.8.4.4
dns_nameserver: {{ google_nameserver }}

pg_host: sql
pg_ip: 192.168.20.22
pg_subnet: 192.168.20.0/24

log_host: sql
log_host_ip: 192.168.20.22
```

Use ansible-vault for secrets

- Use ansible-vault to store secrets safely.

```
$ touch group_vars/dbservers/secrets
```

```
$ ansible-vault encrypt group_vars/dbservers/secrets
```

```
Vault password:
```

```
Confirm Vault password:
```

```
Encryption successful
```

```
$ ansible-vault edit group_vars/dbservers/secrets
```

```
Vault password:
```

```
$ cat group_vars/dbservers/secrets
```

```
$ANSIBLE_VAULT;1.1;AES256
```

```
30623164636337303064313565393361656437343739396235643861336265373138653965303861  
3933306333636164353330393137633061653230366664310a313734323363306261353339306434  
31623732373933333666656665646135656637356366646231336161323838313661636232613365  
6431636132373036300a666633336135376361326163633961626231396433393533663064306336  
65306365323836633838306639336230383039353035343239306432313535326633
```


Variable precedence

- role defaults
- inventory vars
- inventory group_vars
- inventory host_vars
- playbook group_vars
- playbook host_vars
- host facts
- registered vars
- set_facts
- play vars
- play vars_prompt
- play vars_files
- role and include vars
- block vars (only for tasks in block)
- task vars (only for the task)
- extra vars (always win precedence)

ansible_ssh_user

- Don't login as root
- Settle on become_method: sudo/su/doas
- Don't use service account interactively
echo logout > ~/.bash_profile
- Learn how to use signed ssh keys.

```
TrustedUserCAKeys /etc/ssh/ca_key.pub  
AuthorizedKeysFile /dev/null
```

Use this file: ansible.cfg

- current directory
- home directory
- /etc/ansible/
ansible.cfg

Has/hides ansible-playbook options so your commands can be short.

```
[defaults]
hostfile = vagrant.ini
host_key_checking = False
ask_vault_pass = True
retry_files_enabled = False
executable = sh
remote_tmp=
log_path = /tmp/ansible_run.log

gather_facts = smart
roles_path = galaxy_roles:roles

[privilege_escalation]
become_method = sudo

[ssh_connection]
scp_if_ssh = True
pipelining = True
```

Blocks

- Error handling

```
tasks:
- block:
  - debug: msg='i execute normally'
  - command: /bin/false
  - debug: msg='i never execute, cause ERROR!'
rescue:
  - debug: msg='I caught an error'
  - command: /bin/false
  - debug: msg='I also never execute :-( '
always:
  - debug: msg="this always executes"
```

Add tests to your playbooks

- `--tags test`
- fail early.
- keep adding test cases.

```
- name: 'verify listening on port 9200'
  wait_for: port=9200 host={{ inventory_hostname }} timeout=5
  tags:
    - test

- name: 'query cluster health'
  uri: url='http://{{inventory_hostname}}:9200/_cluster/health?pretty'
  register: cluster
  tags:
    - test

- name: 'verify that nodes have joined the cluster'
  assert:
    that:
      - cluster.json.number_of_data_nodes != 0
  tags:
    - test
```

Manage your dependencies

requirements.txt:
extra Python
packages for
modules.

requirements.yml:
list of extra roles to
import.

```
apache-libcloud      # galaxy role
azure                - src: hostclick.tomcat
boto
ConfigParser
cs                   # role in git
django              - src: https://github.com/bbaassssiiee/RHEL6-STIG.git
docker-py           version: devel
dopy
httplib2
linode-python
netaddr
psycpg2
pycrypto
pycurl
pysphere
python-keyczar
python-keystoneclient
python-novaclient
python-quantumclient
pretty
pyrax
requests
shade
wsgiref
zmq
```

Separate top level playbooks by role

- --limit

```
#!/usr/bin/env ansible-playbook
# master playbook
- name: 'site.yml'
  hosts: all:!local
  remote_user: "{{ ansible_user }}"
```

```
- include: dbservers.yml
- include: webservers.yml
```

```
#!/usr/bin/env ansible-playbook
# playbook for dbserver tier
```

```
- name: 'dbservers.yml'
  hosts: dbservers
  become: yes
  gather_facts: True
```

```
roles:
  - common
  - postgres
```

Decouple roles

- Roles are ways of automatically loading certain vars_files, tasks, and handlers based on a known file structure.
- Grouping content by roles also allows easy sharing of roles with other users.
- <http://galaxy.ansible.com>

```
ssh
|-- files
|   |-- bash_profile
|-- handlers
|   |-- main.yml
|-- tasks
|   |-- main.yml
|-- templates
|   |-- ssh_config.j2
|   |-- sshd_config.j2
```


Roles precedence

- Separate your own roles and Galaxy roles.
- Leave room for “corporate roles”.
- Configure `roles_path` to search for roles.
- `galaxy_roles:roles:/etc/ansible/roles`

Organization of a role

- defaults
- files
- **handlers**
- library
- meta
- **tasks/main.yml**
- **templates**
- tests
- vars

```
---  
  
- name: 'ensure package ntp is installed'  
  package: name=ntp state=latest  
  tags:  
    - ntp  
  
- name: 'build and write /etc/ntp.conf file'  
  template: src=ntp.conf.j2 dest=/etc/ntp.conf  
               owner=root group=root mode=0644 backup=no  
  notify:  
    - restart ntpd  
  tags:  
    - ntp  
  
- name: 'enable ntpd as a service'  
  service: name=ntpd enabled=yes state=started  
  tags:  
    - ntp
```

Handlers for a role

- defaults
- files
- **handlers**
- library
- meta
- tasks
- templates
- tests
- vars

```
---
```

```
- name: restart ntpd  
  service: name=ntpd state=restarted
```

Templates for a role

- defaults
- files
- handlers
- library
- meta
- tasks
- **templates**
- tests
- vars

```
tinker panic 0
restrict 127.0.0.1
restrict default kod nomodify notrap
```

```
# NTP server pool, use 3 minimally
{% for server in ntp_servers %}
server {{ server }}
{% endfor %}
```

```
driftfile /var/lib/ntp/drift
```

```
# Avoid using the local clock
server 127.127.1.0 # local clock
fudge 127.127.1.0 stratum 10
```

Default vars for a role

- **defaults/main.yml**
- files
- handlers
- library
- meta
- tasks
- templates
- tests
- vars

```
---  
  
# ntp  
ntp_servers:  
  # use nearby ntp servers if available  
  - 0.pool.ntp.org  
  - 1.pool.ntp.org  
  - 2.pool.ntp.org  
  - 3.pool.ntp.org
```

Bundling modules with a role

- defaults
- files
- handlers
- **library/pam.py**
- meta
- **tasks/main.yml**
- templates
- tests
- vars

```
-  
- name: 'Disable accounts after three login failures in a 15-minute interval.'  
  pam:  
    service: "{{ item }}"  
    type: auth  
    control: required  
    pam_module: pam_faillock.so  
    before_line: auth sufficient pam_unix.so try_first_pass  
    arguments: preauth silent deny=3 unlock_time=604800 fail_interval=900  
    state: present  
  with_items:  
    - password-auth  
    - system-auth  
  tags:  
    - pam
```

Management

- Is the team applying best practice and common principles when building / configuring applications and environments?
- Is the team automating the deployment and promotion process for applications and environments?
- Is the team validating functional and non-functional requirements (automatically) before promoting applications to Production environments?
- Is the team providing and receiving the feedback needed to be in control of the solutions they are responsible for?
- Is the team pro-actively collaborating with other teams to integrate applications and environments and share technical knowledge?

Benchmarks

A good Config Management Strategy should allow to answer all of the following questions with „yes“:

- Can I exactly reproduce any of my environments, including the version of the operating system, its patch level, the network configuration, the SW stack, the application deployed into it, and their configuration?
- Can I easily make an incremental change to any of these individual items and deploy the change to any, and all, of my environments?
- Can I easily see each change that occurred to a particular environment and trace it back to see exactly what the change was, who made it, and when?
- Can I satisfy all of the compliance regulations that I am subject to?
- Is it easy for every member of the team to get the information that they need, and to make the changes that they need to make? Or does the strategy get in the way of efficient delivery, increasing cycle time and decreasing feedback?

Ansible, best practices.

Bas Meijer 2016 Amsterdam